


## HW5-MATH4322


Anthony Castillo-ID:1670011

2022-11-04

Question 1 (a):



Question 1 (b):



Question 2:

```
x <- c(0.1,0.15,0.2,0.2,0.55,0.6,0.6,0.65,0.7,0.75)
mean(x)

## [1] 0.45
```

The majority vote approach will have the final classification be Red because there are six values greater than 0.5 making them red and only 4 values less than 0.5 making them green. The average probability approach will have the final classification of Green since the average is less than 0.5.

Question 3:

1. use recursive binary splitting to grow tree on the training data
2. apply pruning to the tree to obtain the best sequence of best subtrees
3. use k-fold cross validation to prune the tree
4. average the results for each value and choose the best value
5. return subtree from step 2 that corresponds to chosen value

Question 4 (a):

```
library(ISLR2)
n <- nrow(OJ)
Train <- sample(1:n,800)
oj.train <- OJ[Train,]
oj.test <- OJ[-Train,]
```

Question 4 (b):

```

library(tree)
oj.tree <- tree(Purchase~., data = oj.train)
summary(oj.tree)

##
## Classification tree:
## tree(formula = Purchase ~ ., data = oj.train)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "ListPriceDiff"
## Number of terminal nodes: 7
## Residual mean deviance: 0.7765 = 615.7 / 793
## Misclassification error rate: 0.1662 = 133 / 800

```

The training error rate is 0.1962 and it has 8 terminal nodes

Question 4 (c):

```

oj.tree

## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 800 1073.00 CH ( 0.60625 0.39375 )
##    2) LoyalCH < 0.48285 301 328.90 MM ( 0.23588 0.76412 )
##      4) LoyalCH < 0.280875 163 113.30 MM ( 0.11043 0.88957 ) *
##      5) LoyalCH > 0.280875 138 183.80 MM ( 0.38406 0.61594 )
##        10) PriceDiff < -0.24 15 0.00 MM ( 0.00000 1.00000 ) *
##        11) PriceDiff > -0.24 123 168.20 MM ( 0.43089 0.56911 ) *
##    3) LoyalCH > 0.48285 499 455.50 CH ( 0.82966 0.17034 )
##      6) LoyalCH < 0.740621 220 276.70 CH ( 0.67727 0.32273 )
##        12) ListPriceDiff < 0.235 92 126.40 MM ( 0.44565 0.55435 )
##          24) PriceDiff < 0.015 49 56.70 MM ( 0.26531 0.73469 ) *
##          25) PriceDiff > 0.015 43 55.62 CH ( 0.65116 0.34884 ) *
##      13) ListPriceDiff > 0.235 128 111.00 CH ( 0.84375 0.15625 ) *
##      7) LoyalCH > 0.740621 279 111.10 CH ( 0.94982 0.05018 ) *

```


Terminal Node 22 has a split of LoyalCH < 0.277977 and the deviance is 58.63 and is predicted for value MM which has a probability of 71% of being value MM and 29% of being value CH

Question 4 (d):

```

plot(oj.tree)
text(oj.tree)

```



The value CH has a higher chance of being predicted the value MM

Question 4 (e):

```

oj.pred <- predict(oj.tree,newdata = oj.test, type = "class")
table(pred=oj.pred,true=oj.test$Purchase)

##      true
## pred  CH  MM
##   CH 139  17
##   MM  29  85

#test error rate
mean(oj.pred != oj.test$Purchase)

## [1] 0.1703704

```

Question 4 (f):

```

set.seed(10)
oj.cv <- cv.tree(oj.tree)
oj.cv


## $size
## [1] 7 6 5 4 3 2 1
##
## $dev

```

```
## [1] 735.6377 728.0531 703.7670 735.5013 744.5177 803.3631 1074.0156
##
## $k
## [1] -Inf 14.13538 15.66329 31.78660 39.31949 67.72780 288.25708
##
## $method
## [1] "deviance"
##
## attr(,"class")
## [1] "prune" "tree.sequence"
```

Question 4 (g):

```
plot(oj.cv$size,oj.cv$dev,type="b")
```




Question 4 (h):

The lowest cross-validated classification error rate is for tree size = 7

Question 4 (i):

```
oj.prune <- prune.tree(oj.tree, best = 7)
plot(oj.prune)
text(oj.prune)
```



Question 4 (j):

```
summary(oj.prune)

##
## Classification tree:
## tree(formula = Purchase ~ ., data = oj.train)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "ListPriceDiff"
## Number of terminal nodes: 7
## Residual mean deviance: 0.7765 = 615.7 / 793
## Misclassification error rate: 0.1662 = 133 / 800

summary(oj.tree)

##
## Classification tree:
## tree(formula = Purchase ~ ., data = oj.train)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "ListPriceDiff"
## Number of terminal nodes: 7
## Residual mean deviance: 0.7765 = 615.7 / 793
## Misclassification error rate: 0.1662 = 133 / 800
```

The pruned tree has a higher training error rate

Question 4 (k):

```
oj.prune.pred <- predict(oj.prune, newdata = oj.test, type="class")
table(pred=oj.prune.pred,true=oj.test$Purchase)

##      true
## pred  CH  MM
##   CH 139  17
##   MM  29  85

#pruned tree test error rate
mean(oj.prune.pred != oj.test$Purchase)

## [1] 0.1703704
```

The pruned tree has a higher test error rate

Question 5 (a):

```
library(ISLR)


##
## Attaching package: 'ISLR'

## The following objects are masked from 'package:ISLR2':
##
##      Auto, Credit

n <- nrow(Carseats)
Train <- sample(1:n,0.70*n)
carseats.train <- Carseats[Train,]
carseats.test <- Carseats[-Train,]
```

Question 5 (b):

```
library(tree)
carseats.tree <- tree(Sales~., data = carseats.train)
plot(carseats.tree)
text(carseats.tree, pretty = 0)
```



```
summary(carseats.tree)

##
## Regression tree:
## tree(formula = Sales ~ ., data = carseats.train)
## Variables actually used in tree construction:
## [1] "ShelveLoc" "Price" "Age" "Advertising" "Education"
## [6] "Population" "CompPrice"
## Number of terminal nodes: 16
## Residual mean deviance: 2.369 = 625.4 / 264
## Distribution of residuals:
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## -3.6630 -0.9307 -0.0675 0.0000 0.9643 5.1090


#MSE
carseats.pred <- predict(carseats.tree, newdata = carseats.test)
mean((carseats.pred-carseats.test$Sales)^2)

## [1] 4.756139
```

Question 5 (c):

```
set.seed(11)
cv.carseat <- cv.tree(carseats.tree)
plot(cv.carseat$size,cv.carseat$dev, type="b")
```





```
prune.carseat <- prune.tree(carseats.tree,best = 15)
carseat.pred.pruned <- predict(prune.carseat, newdata = carseats.test)
mean((carseat.pred.pruned-carseats.test$Sales)^2)

## [1] 4.793973
```

Yes pruning the tree improved the test MSE  
Question 5 (d):

```
library(randomForest)

## randomForest 4.7-1.1

## Type rfNews() to see new features/changes/bug fixes.

set.seed(10)
car.bag <- randomForest(Sales~., data = carseats.train,mtry = ncol(Carseats)-
1, importance = TRUE)
car.bag

##
## Call:
## randomForest(formula = Sales ~ ., data = carseats.train, mtry = ncol(Cars
eats) - 1, importance = TRUE)
##              Type of random forest: regression
##              Number of trees: 500
## No. of variables tried at each split: 10
```

```
##
##           Mean of squared residuals: 2.483044
##           % Var explained: 72.01

#test MSE
car.bagpred <- predict(car.bag ,newdata = carseats.test)
mean((car.bagpred-carseats.test$Sales)^2)


## [1] 2.580592

importance(car.bag)

##           %IncMSE  IncNodePurity
## CompPrice    23.974062    195.587895
## Income       11.942030    133.716750
## Advertising   25.535646    173.320606
## Population    2.711433     88.995413
## Price        71.662258    689.126045
## ShelfLoc     78.376524    854.486203
## Age          20.279028    202.421144
## Education     6.803589     69.914656
## Urban        -0.561290     10.643511
## US           2.778464      9.561316

varImpPlot(car.bag)
```


car.bag



Price and ShelfLoc are the most important variables  
Question 5 (e):

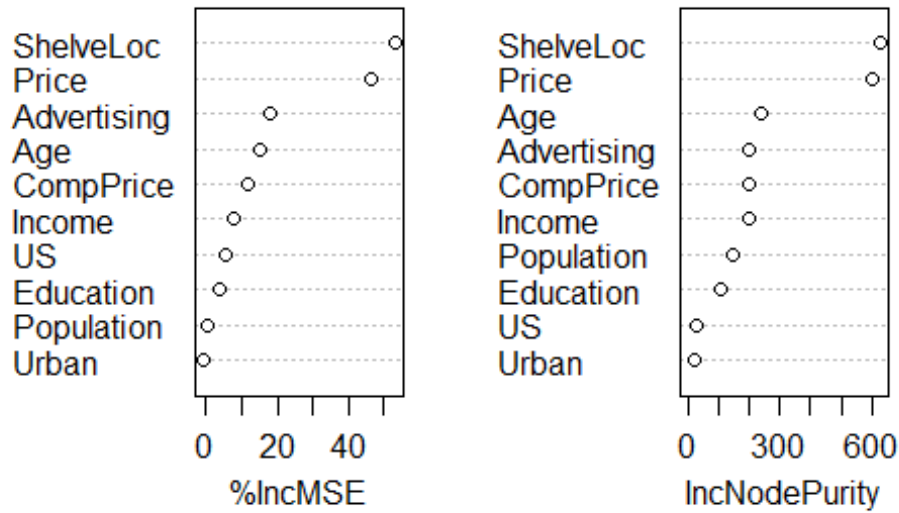
```
for(i in c(1,3,5)){  
  set.seed(10)  
  car.bag <- randomForest(Sales~., data = carseats.train,  
                          mtry = i,  
                          importance = TRUE)  
  #print(importance(car.bag))  
  varImpPlot(car.bag)  
  
  car.bagpred <- predict(car.bag ,newdata = carseats.test)  
  print(mean((car.bagpred-carseats.test$Sales)^2))  
}
```

car.bag



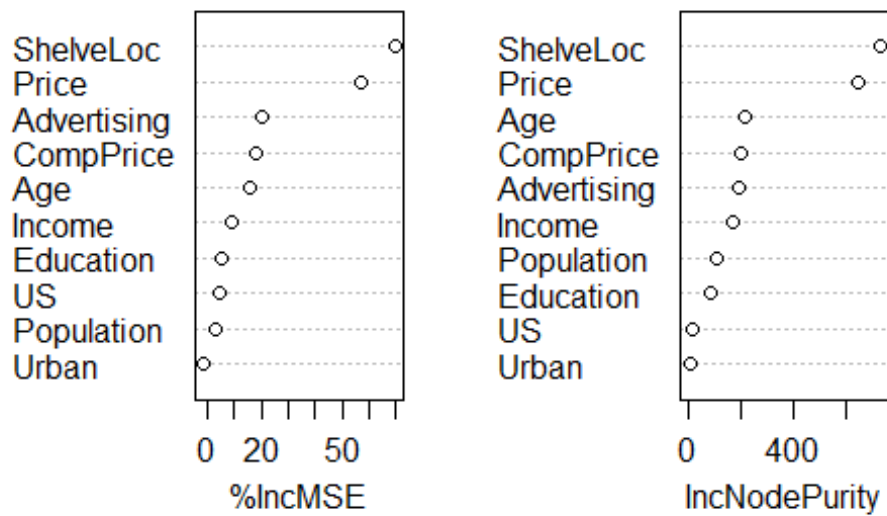
```
## [1] 3.505526
```

### car.bag



```
## [1] 2.471869
```

### car.bag



```
## [1] 2.476211
```

Price and ShelfLoc stay the most important variable throughout all the m's  
The effect of m as it is increased the MSE test error decreases with each m  
Question 6 (a):

```
library(ISLR2)
Train <- sample(1:1000)
Caravan$Purchase = ifelse(Caravan$Purchase == "Yes", 1, 0)
caravan.train <- Caravan[Train,]
caravan.test <- Caravan[-Train,]
```

Question 6 (b):


```
library(gbm)
## Loaded gbm 2.1.8.1

set.seed(1)
caravan.boost <- gbm(Purchase~., data = caravan.train,
                     distribution = "gaussian",
                     n.tree = 1000,
                     shrinkage = 0.01,)

## Warning in gbm.fit(x = x, y = y, offset = offset, distribution = distribut
ion, :
## variable 50: PVRAAUT has no variation.

## Warning in gbm.fit(x = x, y = y, offset = offset, distribution = distribut
ion, :
## variable 71: AVRAAUT has no variation.

summary(caravan.boost)
```



```
##          var    rel.inf
## PPERSAUT PPERSAUT 14.1789945
## MKOOPKLA MKOOPKLA  9.4143251
## MOPLHOOG MOPLHOOG  6.9794184
## MBERMIDD MBERMIDD  5.8329731
## PBRAND    PBRAND   5.1588061
## MGODGE    MGODGE   4.4264998
## MOSTYPE   MOSTYPE  3.7376459
## MINK3045  MINK3045  3.6212597
## ABRAND    ABRAND   3.4895489
## MAUT1     MAUT1    3.2085799
## MAUT2     MAUT2    3.1092005
## MSKA      MSKA     3.0582519
## PWAPART   PWAPART  2.9361318
## MGODPR    MGODPR   2.1424605
## MBERARBG  MBERARBG  1.9768421
## MSKC      MSKC     1.8684321
## MBERHOOG  MBERHOOG  1.8637551
## MINK7512  MINK7512  1.4403762
## PBYSTAND  PBYSTAND  1.4359120
## MINKGEM   MINKGEM   1.4300349
## MGODOV    MGODOV   1.3682377
## MFEKIND   MFEKIND   1.3047131
## MRELOV    MRELOV   1.2748401
## MFGKIND   MFGKIND   1.1939343
## MHHUUR    MHHUUR   1.1805029
## MSKB1     MSKB1    0.9297678
```

|             |          |           |
|-------------|----------|-----------|
| ## MRELGE   | MRELGE   | 0.8651113 |
| ## MGODRK   | MGODRK   | 0.7833458 |
| ## MZPART   | MZPART   | 0.7385279 |
| ## MSKD     | MSKD     | 0.7201284 |
| ## MBERARBO | MBERARBO | 0.7187453 |
| ## APERSAUT | APERSAUT | 0.6886313 |
| ## MOPLMIDD | MOPLMIDD | 0.6739193 |
| ## MGEMOMV  | MGEMOMV  | 0.6377615 |
| ## MINKM30  | MINKM30  | 0.6265648 |
| ## MZFONDS  | MZFONDS  | 0.6196066 |
| ## MAUT0    | MAUT0    | 0.5649730 |
| ## MHKOOP   | MHKOOP   | 0.4374713 |
| ## MOSHOOFD | MOSHOOFD | 0.4315327 |
| ## MINK4575 | MINK4575 | 0.4136190 |
| ## MGEMLEEF | MGEMLEEF | 0.4132185 |
| ## MINK123M | MINK123M | 0.3967414 |
| ## MRELSA   | MRELSA   | 0.3455787 |
| ## MBERBOER | MBERBOER | 0.3095742 |
| ## MSKB2    | MSKB2    | 0.2661275 |
| ## PMOTSCO  | PMOTSCO  | 0.2355198 |
| ## MBERZELF | MBERZELF | 0.2142473 |
| ## MFALLEEN | MFALLEEN | 0.1690279 |
| ## MOPLLAAG | MOPLLAAG | 0.1685822 |
| ## MAANTHUI | MAANTHUI | 0.0000000 |
| ## PWABEDR  | PWABEDR  | 0.0000000 |
| ## PWALAND  | PWALAND  | 0.0000000 |
| ## PBESAUT  | PBESAUT  | 0.0000000 |
| ## PVRAAUT  | PVRAAUT  | 0.0000000 |
| ## PAANHANG | PAANHANG | 0.0000000 |
| ## PTRACTOR | PTRACTOR | 0.0000000 |
| ## PWERKT   | PWERKT   | 0.0000000 |
| ## PBROM    | PBROM    | 0.0000000 |
| ## PLEVEN   | PLEVEN   | 0.0000000 |
| ## PPERSONG | PPERSONG | 0.0000000 |
| ## PGEZONG  | PGEZONG  | 0.0000000 |
| ## PWAOREG  | PWAOREG  | 0.0000000 |
| ## PZEILPL  | PZEILPL  | 0.0000000 |
| ## PPLEZIER | PPLEZIER | 0.0000000 |
| ## PFIETS   | PFIETS   | 0.0000000 |
| ## PINBOED  | PINBOED  | 0.0000000 |
| ## AWAPART  | AWAPART  | 0.0000000 |
| ## AWABEDR  | AWABEDR  | 0.0000000 |
| ## AWALAND  | AWALAND  | 0.0000000 |
| ## ABESAUT  | ABESAUT  | 0.0000000 |
| ## AMOTSCO  | AMOTSCO  | 0.0000000 |
| ## AVRAAUT  | AVRAAUT  | 0.0000000 |
| ## AAANHANG | AAANHANG | 0.0000000 |
| ## ATRACTOR | ATRACTOR | 0.0000000 |
| ## AWERKT   | AWERKT   | 0.0000000 |
| ## ABROM    | ABROM    | 0.0000000 |

```
## ALEVEN      ALEVEN  0.0000000
## APERSONG    APERSONG 0.0000000
## AGEZONG     AGEZONG  0.0000000
## AWAOREG     AWAOREG  0.0000000
## AZEILPL     AZEILPL  0.0000000
## APLEZIER    APLEZIER 0.0000000
## AFIETS      AFIETS   0.0000000
## AINBOED     AINBOED  0.0000000
## ABYSTAND    ABYSTAND 0.0000000
```

Question 6 (c):

```
set.seed(10)
pred.carboost <- predict(caravan.boost,
                          newdata = caravan.test,
                          n.trees = 1000)
greater20 <- ifelse(pred.carboost > .20,1,0)
table(pred = greater20,true= caravan.test$Purchase)

##      true
## pred    0    1
##    0 4501  279
##    1   32   10

#predicted to make purchase and in fact make one
print(10/(10+279))

## [1] 0.03460208
```